



Naive Mult vs. cuBLAS for Different Matrix Sizes



*Linear graphs do not capture changes at small values well, which is why log(GFLOPS) is used


—●— Naïve Kernel GFLOPS —●— cuBLAS GFLOPS

Tiled Mult vs. cuBLAS for Different Matrix Sizes



—●— Kernel GFLOPS —●— cuBLAS GFLOPS

Vectorized 2d Coarse Mult vs. cuBLAS for Different Matrix Sizes



—●— Kernel GFLOPS —●— cuBLAS GFLOPS